

Conformal Prediction for Reliable Handover Under Distribution Shift

Johan Eliasson Eklund
<https://github.com/elitan>

Abstract—ML handover prediction is accurate in-distribution but fragile under shift. We evaluate conformal prediction (CP) for 5G handover under synthetic and real-world drift. We compare static CP, Adaptive Conformal Inference (ACI), dynamic-step ACI (Daci), weighted CP, and a confidence-gated triggered ACI against Top- k and a 3dB hysteresis baseline. On in-distribution synthetic data, static CP reaches 90.2% coverage with set size 2.45. Under speed and noise shifts, weighted CP reaches 91.1% in the speed-shift case. Under severe shadow shift, static CP drops to 69.2% while ACI restores 88.8% and Daci reaches 89.9% at higher overhead. In regime-switch streams, ACI stabilizes rolling coverage near the 90% target, while Daci reaches 91.8%. ACI step-size sweep shows $\gamma = 0.002\text{--}0.005$ gives the best reliability-overhead tradeoff in our setup. On Irish 5G driving traces with speed-split drift, static CP reaches 76.4% coverage (size 4.69), weighted CP 73.6% (size 4.22), ACI 87.9% (size 14.7), and Daci 92.5% (size 15.9). Results show reliability under shift needs adaptive conformal control, not static calibration alone.

I. INTRODUCTION

Predictive handover reduces latency but can fail badly when radio conditions shift. Traditional 3GPP handover logic relies on hysteresis events [2], [3]. ML-based handover prediction improves point accuracy but provides no risk control when distribution changes [4], [12]. In production, that gap maps directly to radio link failure risk.

Conformal prediction (CP) gives finite-sample marginal coverage guarantees [5], [6]. Recent wireless CP work focuses on demodulation, channel tasks, and beam selection [7]–[9], while handover under distribution shift is still underexplored. This paper targets that gap.

Contributions.

- 1) We benchmark handover reliability under four synthetic shifts plus a regime-switch stream.
- 2) We compare static CP, ACI [10], dynamic-step ACI, weighted CP, and confidence-gated triggered ACI under the same base predictor and KPI mapping.
- 3) We validate on Irish 5G driving traces with source-target speed split.
- 4) We provide budget-aware reproducible runs (local-first, capped overflow policy) and release all generated artifacts.

II. RELATED WORK

ML handover methods span supervised and reinforcement-learning policies [4], [12]. CP in wireless has shown value for calibration and reliability [9]. CP for beam selection shows strong reliability-efficiency tradeoffs [7], [8]. The missing

piece is handover under shift: static calibration can fail as mobility, shadowing, and measurement noise drift over time.

III. SYSTEM AND METHODS

A. Handover Prediction Setup

At time t , the model predicts future best cell $y_t = \arg \max_k \text{RSRP}_k(t+H)$ using input

$$\mathbf{x}_t = [\text{RSRP}_1, \dots, \text{RSRP}_K, \mathbf{e}_{ct}, v_t]. \quad (1)$$

We use an MLP classifier and softmax scores $\hat{p}(y | \mathbf{x})$.

B. Baselines and Conformal Variants

3dB baseline: handover if best neighbor exceeds serving by 3dB.

Static CP:

$$\mathcal{C}(\mathbf{x}) = \{y : \hat{p}(y | \mathbf{x}) \geq 1 - \hat{q}\}, \quad (2)$$

with \hat{q} calibrated on held-out source calibration data.

ACI: online update of effective miscoverage level to track sequential drift [10].

Daci: dynamic-step ACI that switches between low/high update rates using an EMA of recent errors.

Weighted CP: source calibration scores reweighted by estimated density ratio $w(\mathbf{x}) \propto p_T(\mathbf{x})/p_S(\mathbf{x})$ using a source-vs-target logistic discriminator.

Triggered ACI: confidence-gated mixture that uses static CP on high-confidence samples and ACI sets on low-confidence samples (threshold from source calibration confidence quantile).

C. System KPI Mapping

Coverage maps to handover success with bounded miss risk. Set size maps to measurement overhead. Undercoverage maps to RLF proxy. Serving-cell retention in small sets acts as implicit hysteresis and affects ping-pong rate.

IV. EXPERIMENTAL SETUP

A. Synthetic Shift Benchmark

Source setting: medium scenario (4×4 cells, $\sigma = 6$ dB shadowing, measurement noise 4 dB, speed 1–30 m/s, horizon $H = 10$). We train on source and calibrate on source only.

Target shifts:

- 1) IID (same as source)
- 2) Speed shift (20–50 m/s)
- 3) Measurement-noise shift (8 dB)

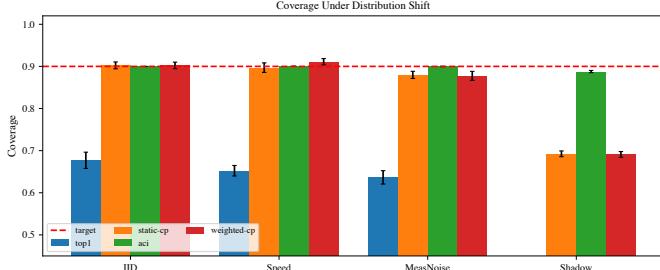


Fig. 1. Coverage across synthetic shifts. Static CP degrades under hard shift, ACI is most robust, weighted CP helps in moderate covariate shift.

TABLE I
SYNTHETIC SHIFT COVERAGE (MEAN \pm STD)

Shift	3dB	Top-1	Top-3	Static CP	ACI
IID	.64 \pm .02	.68 \pm .02	.91 \pm .01	.90 \pm .01	.90 \pm .00
Speed	.62 \pm .01	.65 \pm .02	.90 \pm .01	.90 \pm .01	.90 \pm .00
MeasNoise	.55 \pm .02	.64 \pm .02	.89 \pm .01	.88 \pm .01	.90 \pm .00
Shadow	.40 \pm .01	.43 \pm .01	.72 \pm .01	.69 \pm .01	.89 \pm .00
Regime	.49 \pm .02	.53 \pm .02	.79 \pm .01	.78 \pm .02	.89 \pm .00

- 4) Shadow shift ($\sigma = 10$ dB)
- 5) Regime switch (source-like first half, harsh second half)

Each result is mean \pm std across 5 seeds (42,123,456,789,1011), 600 trajectories/seed, 20 epochs.

B. Real-World Drift Benchmark

Dataset: Irish 5G driving traces [13]. We split traces by average speed: lower-speed traces as source, higher-speed traces as target. Model trains and calibrates on source only, then evaluates on target.

V. RESULTS

A. Coverage Under Shift

Table I shows the main trend: static CP is reliable near source but degrades under strong shift (shadow, regime). ACI keeps coverage close to target by expanding sets online.

B. Tradeoff in Hard Shifts

Paired seed deltas confirm hard-shift reliability gains: ACI vs static is +19.5pp coverage on shadow shift (95% CI [18.9, 20.1]) and +11.1pp on regime switch (95% CI [10.0, 12.2]). DACI further improves over static by +20.7pp (shadow) and +13.6pp (regime), and over ACI by +1.1pp and +2.6pp, with overhead increases of +4.81pp and +6.01pp, respectively.

C. Regime-Switch Stability

In regime-switch streams, static and weighted thresholds lag after the phase boundary. ACI adapts online and recovers target-level coverage.

TABLE II
HARD-SHIFT KPI TRADEOFF (MEAN OVER 5 SEEDS)

Shift	Method	Coverage	Set Size	RLF Proxy	Overhead
Shadow	Static CP	.692	2.62	.306	.170
	ACI	.888	5.98	.072	.601
	DACI	.899	6.35	.058	.650
	Triggered ACI	.842	5.38	.122	.541
	Weighted CP	.691	2.61	.306	.169
Regime	Static CP	.781	2.61	.217	.168
	ACI	.892	4.50	.083	.417
	DACI	.918	5.04	.056	.477
	Triggered ACI	.869	4.22	.108	.387
	Weighted CP	.789	2.70	.209	.176

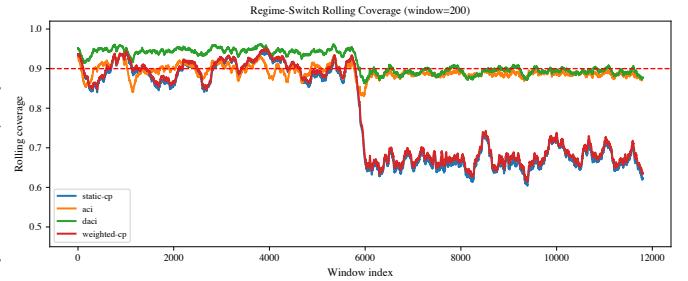


Fig. 2. Rolling coverage in regime-switch stream (window=200). ACI tracks the 90% target more closely than static and weighted CP.

D. ACI Step-Size Sensitivity

Figure 3 quantifies ACI sensitivity. In our regime-switch benchmark, small step sizes ($\gamma = 0.002\text{--}0.005$) achieve the highest coverage ($\approx 89.6\%$) with moderate set inflation, while larger values ($\gamma = 0.05$) reduce coverage to 87.2% but lower overhead. This supports tuning γ as a direct reliability-overhead control knob. Triggered ACI has a second control knob. A regime-switch trigger-quantile sweep in Appendix Fig. A1 shows quantile 0.5 \rightarrow 0.9 improves coverage from 85.5% to 88.3% while overhead rises from 0.358 to 0.409.

E. Real-World Drift Results

Trace-bootstrap on Irish confirms the reliability-cost pattern: ACI vs static gives +11.6pp coverage (95% CI [5.4, 19.4]) with +10.02 set size (95% CI [3.59, 18.84]). DACI vs static gives +16.1pp (95% CI [10.1, 23.7]) with +11.26 set size (95% CI [4.65, 20.17]), and DACI vs ACI adds +4.6pp coverage (95% CI [4.2, 4.9]) with +1.24 set size (95% CI [0.87, 1.70]). A speed-bin breakdown shows high-speed traces are hardest (static 68.9%, weighted 66.4%) while DACI is highest at 88.5% (ACI 83.7%, triggered 77.7%).

VI. DISCUSSION

When to use which method. Static CP is a strong default in stable conditions. Weighted CP helps moderate covariate shift when target feature support overlaps source. Triggered ACI is a middle point when overhead budget is tight. ACI is robust

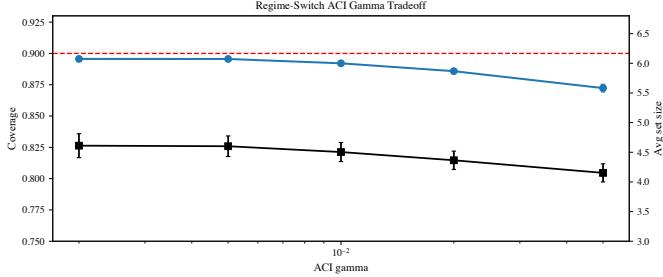


Fig. 3. Regime-switch ACI tradeoff over γ . Smaller γ improves coverage but increases set size and overhead.

TABLE III
IRISH 5G SPEED-SPLIT DRIFT RESULTS

Method	Coverage	Avg Size
Top-1	.307	1.00
Top-3	.625	3.00
Static CP	.764	4.69
ACI	.879	14.70
DACI	.925	15.94
Triggered ACI	.824	10.15
Weighted CP	.736	4.22

for severe sequential drift. DACI is the max-reliability mode when higher overhead is acceptable.

System implications. Reliability gains translate to lower RLF proxy but require explicit overhead budgeting. Table IV gives a simple control policy: low-budget mode prefers static or weighted CP, medium-budget mode uses triggered ACI, high-budget mode uses ACI, and max-reliability mode uses DACI.

Limitations. Synthetic channels still simplify real deployments. Irish traces have limited feature richness versus full network measurement reports. We evaluate offline; online deployment requires streaming integration and control-plane constraints.

VII. CONCLUSION

We presented a shift-focused handover reliability study with conformal prediction. Static CP works well in-distribution but degrades in severe shift. ACI restores near-target coverage under shadow and regime-switch drift. DACI pushes reliability further in hard shifts and on Irish traces, with additional overhead. Weighted CP improves moderate shifts with smaller set inflation than adaptive variants. A confidence-gated triggered ACI recovers a strong middle tradeoff. The core practical result is clear: robust handover reliability needs adaptive conformal control, not static calibration alone.

APPENDIX A APPENDIX: ENSEMBLE AND LATENCY

Measured medium-scenario latency: calibration 0.08 ms, CP set construction 1.59 μ s/sample, NN inference 0.50 μ s/sample.

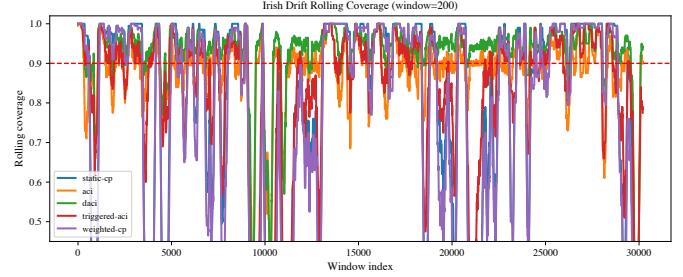


Fig. 4. Irish drift rolling coverage (window=200). DACI reaches highest reliability under source-target speed split.

TABLE IV
BEST METHOD UNDER OVERHEAD CAPS (SYNTHETIC HARD SHIFTS)

Overhead Cap	Shadow Shift	Regime Switch
$\leq .20$	Static CP (.692)	Weighted CP (.789)
$\leq .40$	Static CP (.692)	Triggered ACI (.869)
$\leq .60$	Triggered ACI (.842)	DACI (.918)
$\leq .70$	DACI (.899)	DACI (.918)

REFERENCES

- [1] T. S. Rappaport *et al.*, “Millimeter wave mobile communications for 5G cellular,” *IEEE Access*, 2013.
- [2] 3GPP TS 38.214, “NR; Physical layer procedures for data,” 2022.
- [3] 3GPP TS 38.331, “NR; Radio Resource Control (RRC); Protocol specification,” 2022.
- [4] V. Yajnanarayana, H. Rydén, and L. Hévizi, “5G handover using reinforcement learning,” in *Proc. IEEE 5GWF*, 2020.
- [5] V. Vovk, A. Gammerman, and G. Shafer, *Algorithmic Learning in a Random World*. Springer, 2005.
- [6] A. N. Angelopoulos and S. Bates, “A gentle introduction to conformal prediction,” *arXiv:2107.07511*, 2021.
- [7] D. N. Hegde *et al.*, “Reliable and efficient beam selection using conformal prediction in 6G systems,” in *Proc. IEEE WCM*, 2024.
- [8] J. Deng *et al.*, “SCAN-BEST: Sub-6GHz-aided near-field beam selection with formal reliability guarantees,” *arXiv:2503.13801*, 2025.
- [9] K. M. Cohen *et al.*, “Calibrating AI models for wireless communications via conformal prediction,” *IEEE TMLCN*, 2022.
- [10] I. Gibbs and E. Candès, “Adaptive conformal inference under distribution shift,” *NeurIPS*, 2021.
- [11] Y. Romano *et al.*, “With malice toward none: Assessing uncertainty via equalized coverage,” *HDSR*, 2020.
- [12] W. Lee *et al.*, “Prediction-based conditional handover for 5G mm-wave networks,” *IEEE Access*, 2020.
- [13] D. Raca *et al.*, “Beyond throughput: A 5G dataset with channel and context metrics,” in *Proc. MMSys*, 2020.
- [14] B. Lakshminarayanan *et al.*, “Simple and scalable predictive uncertainty estimation using deep ensembles,” in *NeurIPS*, 2017.

TABLE A1
APPENDIX: CP VS ENSEMBLE (FROM v5 RUNS, 5 SEEDS)

Scenario	CP (1 model)		Ensemble (5 models)	
	Coverage	Size	Coverage	Size
Easy	.900	1.37	.900	1.38
Medium	.893	2.47	.898	2.57
Hard	.898	4.80	.903	5.05

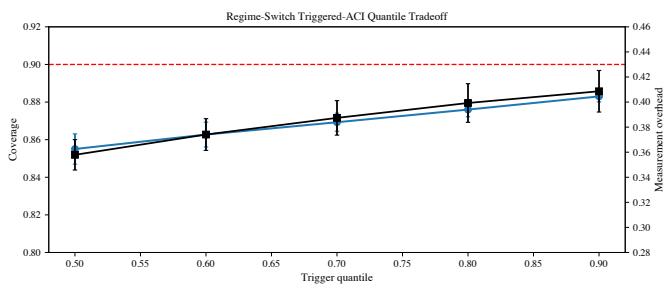


Fig. A1. Regime-switch triggered-ACI tradeoff over confidence-gate quantile.
Higher quantiles increase coverage and overhead.